Transformer Inference Latency on CPU

N = 200 sequential inference calls, single-threaded CPU execution

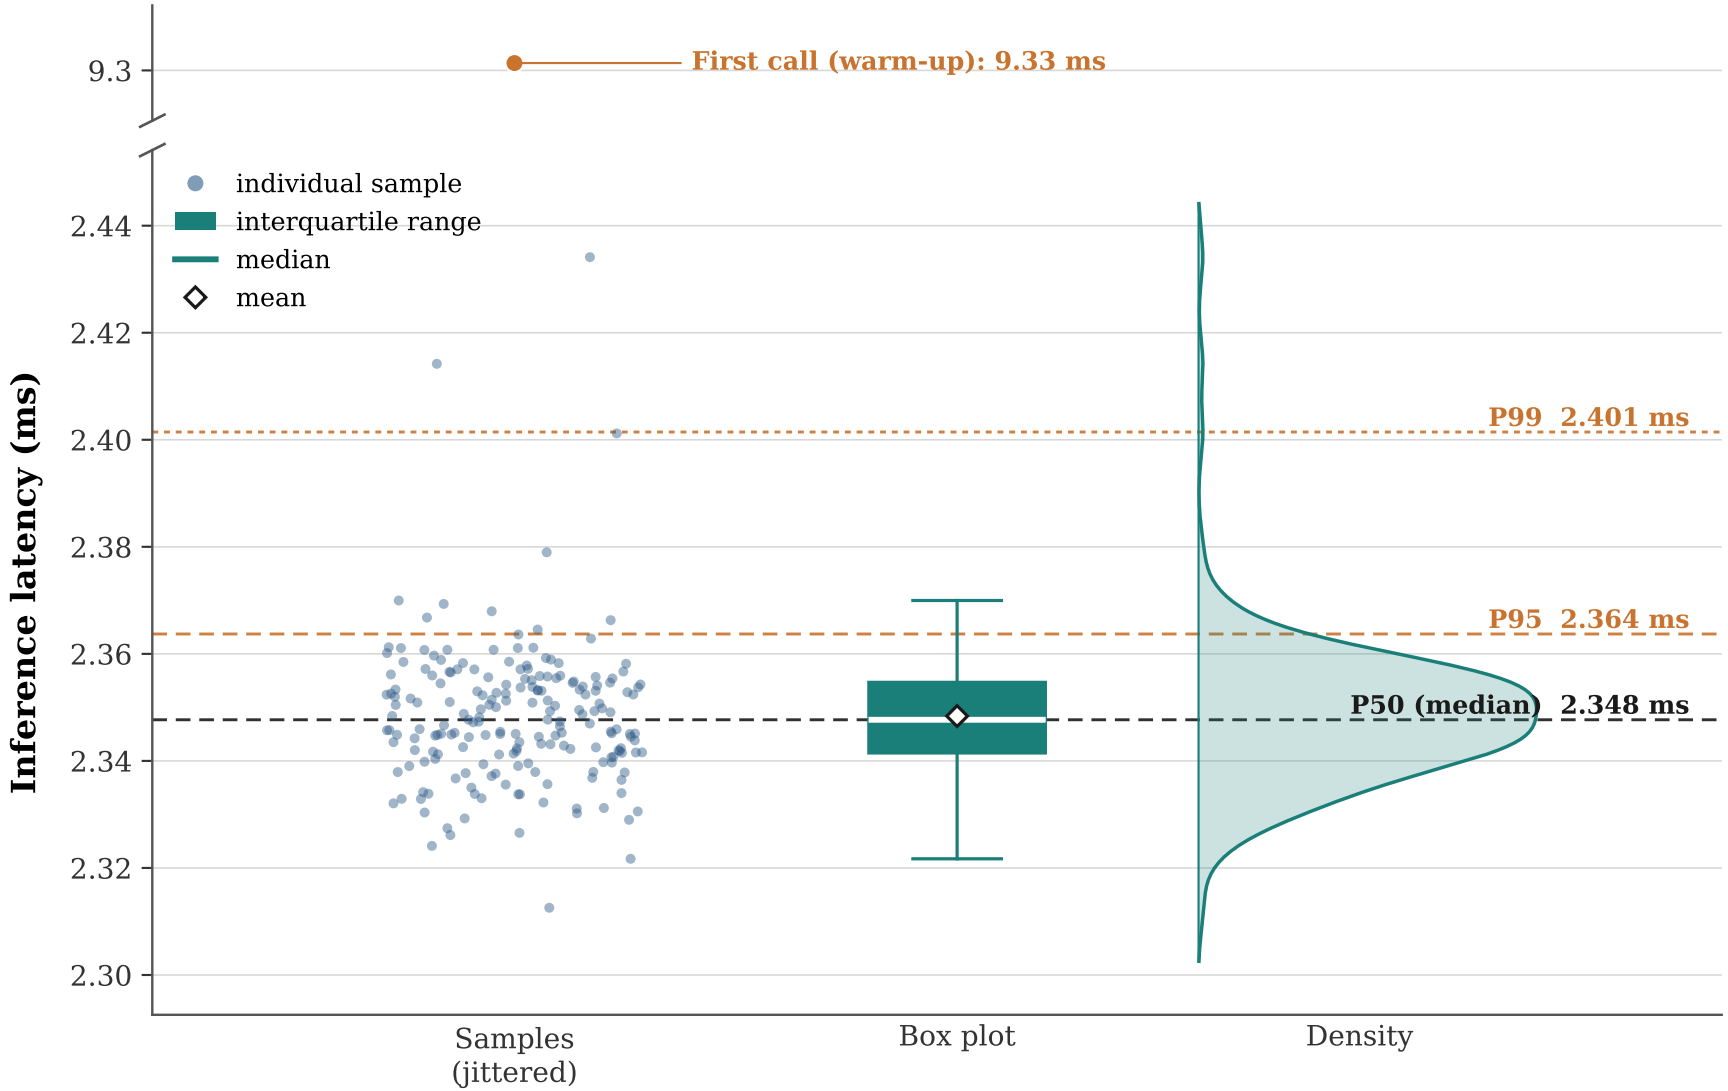

Steady state (N = 199, excl. warm-up): mean 2.348 ± 0.013 ms · median 2.348 ms · P95 2.364 ms · P99 2.401 ms · range 2.313–2.434 ms

All calls (N = 200, incl. warm-up): mean 2.383 ± 0.494 ms · median 2.348 ms